

Study posture for the Foundation exam from day one

Claude Certified Developer, Foundation (CCDV-F) · UST masterclass

EN Session 1 of 12 · English notes

Session 1 did not walk the Pearson booklet, but Vijay already told the room how to *think* so the later exam format is not a surprise. Foundation is a **developer certification**: you will not open an IDE in the test centre and ship an app. You will choose products, name trade-offs, and write how you would implement a scenario. Start that posture now—while Chat and Cowork are still the only screens you have seen.

What you are actually being certified for

This masterclass is CCDV-F, not the business Associate track and not Architect Professional. Prompt engineering still shows up everywhere (Vijay later quoted ~18–20% of items across certs). From day one, practise *explaining* a workflow in Anthropic’s words: client versus model, tokens versus usage meters, Chat versus Cowork versus Code, tool versus MCP versus RAG.

Exam element (as taught in week 1)	What to expect	Study implication from session 1
Item types	MCQ plus written / descriptive scenarios — not application coding in the exam harness	Practise writing “I would use X because...” not pasting a React app
Scenarios	About six use cases you must know thoroughly (detailed on 5 August)	Map each classroom demo (invoices, weather+image, Chat cutoff) to a named pattern
Duration / count (developer track, 6 August slides)	60 questions, 120 minutes (associate quoted the same shape)	Pace: two minutes average; leave time for the long written items
Score (programme guide, 6 August)	Scale described as 1000 marks, pass around 750	You cannot guess your way through the scenario block
Validity / price (classroom)	12 months; associate quoted USD 99, developer USD 125 — confirm on Skilljar	Budget and recertify; numbers move
Delivery	Book your own slot; Pearson mentioned for ID check	Legal name on Skilljar must match photo ID — raise with the org account holder early

NUMBERS ARE CLASSROOM SNAPSHOTS

Prices, pass marks, and question counts came from Vijay’s slides and a July 2026 programme PDF he promised to share. On exam week, the booking page on anthropic.skilljar.com wins.

Where to study (and where not to)

- **Do not default to Udemy.** Vijay’s answer to “is there a Udemy practice set?” was no: use Anthropic Academy on Skilljar.

- **Percipio / internal LMS** (clarified 6 August) is UST’s learning-journey tracker. Completing it is not the Anthropic credential.
- **docs.claude.com** is the glossary the exam will use (including product jargon such as “harness” in a Claude sense, not only ML temperature).
- Internal dumps and unofficial PDFs were mentioned in passing; treat them as unofficial. Prefer the programme guide Anthropic published and the Academy courses (on the order of three to four courses per exam path; ~19 courses visible in the portal in August 2026).

Habits that already showed up in session 1 Q&A

These are exam-shaped habits, not politeness tips.

1. **New chat when the task is done.** Turn three of a thread resends Q1+A1+Q2+A2+Q3. That is tokens and extra hallucination surface.
2. **Write definite prompts:** role, task, input, output format, constraints (Lalit’s five-part list, which Vijay accepted).
3. **Skills are not free.** A short user prompt plus a 1,000-line skill is billed as the concatenation. Reuse locally; do not assume the skill is cached as a discounted prefix unless you are in an API prompt-cache design.
4. **CLAUDE.md is a product feature** Cowork/Code create it—not a one-off file Vijay typed. You will customise it; stuffing entire vendor docs into it is the wrong optimisation.
5. **Estimate cost with evals** when a client asks ROI: parameterise tokens in and out over N interactions. Session 1 already pointed at evaluation rather than a magic calculator in Chat.
6. **QA path:** connectors + generated Playwright (or similar) in CI, not “ignore the course because I am QA.”

Confidentiality and connectors as exam topics

Responsible AI and enterprise governance are on the later syllabus and already appeared: do not enable personal Gmail connectors on org data; attached Cowork folders can be edited or deleted; credentials for connectors sit on the local client; Anthropic’s SLA/compliance story (PII, PCI, GDPR as relevant to the customer geography) is why the *model* can be described as non-retaining while the *chat history* still exists. If an exam stem says “overnight batch, cost-sensitive, not urgent,” you will later recognise Messages Batches versus naive parallel Messages—session 3 previewed that style of item. For now, practise reading stems for *constraint words*: cost, latency, PII, local disk, real-time.

Token literacy without panicking

Enterprise licences are pooled. Hitting a personal “empty tank” is uncommon on day one; wasting a 36-file Cowork sort that a script should own *will* show up in team usage. Session 1’s invoice demo was the right tool for mixed CSV/txt/images that a brittle script cannot parse; it is the wrong tool if you already have a stable schema and a cron job. That judgement—product versus SDK versus RAG—is the Foundation skill.

Your week-1 checklist

- Create or confirm Skilljar access with the name that matches your ID document.
- Bookmark docs.claude.com and the models page; compare them to whatever the desktop picker shows.
- Sit Chat: cutoff question, web search, new thread.
- If Cowork is enabled, attach a *copy* of a folder only—never the only copy of important files.
- Write five sentences: Chat vs Cowork vs Code; client vs LLM; cutoff vs tools; files vs Project RAG; tokens vs the five-hour usage meter (the last two get a full treatment in modules 02–03).

FORWARD POINTER

Module 02 records exam format and Academy/partner counting in detail. Module 03 records the four post-13-July tracks. You do not need to memorise partner Gold/Premier numbers for CCDV-F, but you do need Anthropic’s product vocabulary.

Official documentation

- **Anthropic Academy (Skilljar)**

<https://anthropic.skilljar.com>

Enrolment, self-paced paths, and exam registration as taught in class.

- **Prompt engineering overview**

<https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/overview>

Role, constraints, and output control — the day-one prompt habit, written properly.

- **Claude Docs home**

<https://docs.claude.com/en/home>

Glossary the exam will prefer over informal classroom ASR.

- **Rate limits**

<https://docs.claude.com/en/api/rate-limits>

API-side limits are not the same as Chat’s five-hour usage bar — start noticing the difference.

- **Usage policy**

<https://www.anthropic.com/legal/usage-policy>

What you may send; pairs with the classroom confidentiality rule.

Study notes distilled from the UST CCD Foundation masterclass. Prefer Anthropic’s current docs over session anecdotes. Not official exam material.