

Cowork in practice: folders, connectors, sandbox, and usage meters

Claude Certified Developer, Foundation (CCDV-F) · UST masterclass

EN

Session 2 of 12 · English notes

Cowork is the local task manager Vijay had previewed on 3 August with the invoice pile. On 5 August he slowed down on **consent, connectors, where files actually live, and the usage meter that is not a context window**. If your tenant has disabled Cowork, still learn the model: the exam can ask what the product is for even if you cannot click it today.

Attach a folder like it can delete your work

When you grant a folder, the dialog warns that Claude may **create, edit, delete, and share files with connected tools**. If Gmail is connected, content can be pushed to email. In a locked-down UST tenant you should only see approved connectors; if a personal Gmail connector appears, do not use it on project data. Prefer a **copy** of the folder. Session 1 already showed skip-all-approvals versus manual approve per edit/delete.

Jobs that fit Cowork

Local data means whatever sits on disk: CSV, PPT, PDF, images. Natural language over a pile beats opening every file. Role-shaped examples he gave:

- **Architects:** read SRS, BRD, SOW and draft HLD-style documentation.
- **UI/UX:** from a logo, produce a brand board (colour, type, responsive asset sizes).
- **Developers:** plans, learning notes, code-adjacent docs (heavy coding still belongs in Claude Code).
- **Analysts:** pull stories/epics from Confluence/Jira/Notion (or Google Drive in the demo) and assemble a pre-dev brief.
- **Manual testers:** repeat walkthroughs against a site (not a substitute for Playwright/Selenium—that is Code).
- **Personal:** scheduled blog/LinkedIn drafts; email summaries triggered from the phone if the product still offers that.

Live prompt: read every document in the folder; emit a table (name, created date, key names, amounts, short description). The client showed a **plan/progress rail** (count files → extract → summarise → write spreadsheet). Output can be downloaded; a Google Drive button appeared because *his* demo identity was a Google workspace user—not because files had already left the laptop.

Sandbox: local, managed by Claude

Rajesh asked whether the generated spreadsheet lived in the cloud before download. Vijay: processing is on the local machine inside a **sandbox** Claude manages. You need disk headroom. The path is not a friendly user folder; you can ask Claude where it stored a file. This is the privacy pitch from session 1: proprietary files

are rag'd locally and snippets go to the LLM; the folder is not copied to Anthropic as a file share. Still follow the SLA and never put client secrets in a demo account.

Connectors you actually saw

Customization → connectors. Catalogue items (weather, Apify scrapers, Gmail, Drive, ...) often require **admin allow-list** before import. Connecting Gmail walked OAuth scopes: read mail, manage drafts/send. Be specific in the prompt ("from Gmail") if several mail connectors exist. A later prompt combined Gmail + Drive: unread mail about Project XYZ + latest Drive doc → draft follow-up. Permissions can be allow once / always / deny per tool call. Connector tokens stay on the local client in his telling; compliance (PII, PCI, GDPR as applicable) is why you can trust the vendor's "we do not train on this / we do not keep credentials on the server" claims *if* the product is certified for those regimes—verify for the customer's country.

OUTAGES AND MODEL FALLBACK

During the Gmail demo Sonnet appeared down; the room reported Haiku/4.6-class models still working. Switching model is a valid operational move. Docs will not screenshot that incident.

Usage meters versus tokens on a task

Cowork does not print "this table cost 12,432 tokens" on the canvas. Settings → Usage shows a **percentage of a rolling window**. Classroom sketch: a window on the order of **five hours** that resets so you cannot hammer the LLM on a cheap subscription, plus a **weekly** cap. He illustrated ~1M tokens / five hours and ~10M / week on *his* plan—then immediately said enterprise numbers are admin-configured and not published as a single public figure. End users may not see rupee cost; admins see billing. **API billing is a different configuration**. Heavy image jobs can burn 100% of a window in an hour; then you wait for reset unless the org raises quota.

When Cowork is the wrong spend

Session 1 already argued: a one-off mixed-format invoice sort is cheaper in human time on Cowork than writing a parser. A daily identical job should become a script or an SDK/RAG pipeline. Cowork may emit bash along the way—that is a hint to productise. Token KPI will be a management question; your answer is product choice, not "Claude is expensive."

Memory, schedule, Design beta

Projects can carry instructions, a memory toggle, and schedules. Claude Design (beta in class) sat as a skill-like generator for slides and wireframes. Beta will not pass most enterprise allow-lists—same story as the Chrome extension in the next chapter.

Official documentation

- **Claude desktop / Cowork support**

<https://support.claude.com/>

Where desktop behaviour and availability are documented as they change.

- **MCP**

<https://docs.claude.com/en/docs/agents-and-tools/mcp>

What “connectors” are under the hood.

- **Rate limits (API)**

<https://docs.claude.com/en/api/rate-limits>

Do not confuse with Cowork’s five-hour usage bar.

- **Usage policy**

<https://www.anthropic.com/legal/usage-policy>

What you may send through connectors.

- **Workspaces**

<https://docs.claude.com/en/docs/build-with-claude/workspaces>

Project / workspace grouping in current docs.

Study notes distilled from the UST CCD Foundation masterclass. Prefer Anthropic’s current docs over session anecdotes. Not official exam material.