

Exam prep guide and mock practice

Claude Certified Developer, Foundation (CCDV-F) · UST masterclass

EN Session 12 of 12 · English notes

The last block of the Foundation masterclass was **how to sit**. Vijay had promised this in session 10; on 27 August he walked Skilljar, unofficial mocks, timing, and first-attempt psychology. Treat Skilljar and docs.claude.com as canonical. Treat third-party quiz sites as stamina trainers whose items are *not* the real bank.

Where official content lives

He said most of his teaching was drawn from **Anthropic Academy on Skilljar**. One Foundation-oriented course plus targeted extras beat watching every “AI Fluency” video. He warned: not all Fluency tracks are for this exam — **AI Fluency for builders** was the prompting slice he cared about.

BLUEPRINT PERCENTAGES MOVE

Figures below are what he read from Skilljar on 27 August 2026. Re-read the live exam page before you sit. UI labels (AI Fluency vs prompt engineering) had already been renamed once in the recording.

Weightings he quoted (verify live)

Skilljar / topic cluster (his labels)	Share he quoted	What to practise from this masterclass
Building with the Claude API	~33%	Messages, roles, streaming vs not, models, cost, tools handshake
Tools and MCP (dedicated course if you are weak)	~10% (plus related API items)	Schemas, tool_result, MCP client/server, Cowork connectors
Prompt / AI Fluency for builders	~11% (he summed fluency+tools+context toward ~40% with other slices)	Stop sequences, JSON, templates, extended thinking, cache
Remainder	Applications, agents, RAG, RAI as they appear in the live outline	Modules 10–12 plus product sessions 1–6

His arithmetic was a coaching heuristic (“cover the big course first, then tools, then builder fluency”), not a substitute for the official domain list on exam day.

Mocks: useful, unofficial, repeatable

He dropped two website links in chat (also promised in the repo). He had *not* sat those products himself; he trusted reviews that the **first** was closer to real scenario style. Caveats he actually said:

- Free account; take them more than once; multiple papers (he mentioned four to five sets).
- Each timed sitting can take **120 minutes** — budget a week if you work through several sets properly.

- He mentioned practising “53 or 23” questions in that window depending on the paper. Confirm the live Foundation format (earlier modules: MCQ plus written scenarios). Do not freeze those two numbers as law.
- About **60–80%** similar in feel; straightforward items get twisted; **concepts ~95% the same, questions not the same.**
- High mock scores do not guarantee a pass; they train pace and scenario reading.

Proctor tactics

1. Build a personal pace: how fast you read, analyse, click.
2. Answer immediately when you are sure; **mark for review** when you are not. The proctor lets you skip and return.
3. English of the stem matters more than a keyword hunt — scenario meaning, not a single buzzword.
4. First-attempt pressure: he observed people lose motivation after a fail. Prepare as if this sitting is the one.

Validity and why recertify

Asked about expiry, he said **two years**, then retake. He compared GitHub Copilot credentials going stale in a year of product change. The certification is a snapshot; tools will have moved. Recheck Anthropic’s current credential policy when you register — do not rely on a classroom sentence alone.

A revision map for the last 72 hours

- API: user/assistant/system, stop reasons, streaming vs non-streaming, model families on the live models page.
- Tools: description quality, parallel vs sequential tool_use, tool_result on a user turn, Agent SDK hides the loop.
- MCP: N×N problem, client built into Claude Code, vendor vs your server, mcp.json is location not business logic.
- RAG: retrieve-augment-generate; MongoDB Atlas + Voyage was a **demo**; Bedrock/Chroma/FAISS/Postgres still count.
- Agents: Claude thinks, Claude Code does, SDK loops; built-ins Bash/Monitor/Glob/Grep/WebSearch/WebFetch/AskUser.
- RAI: deny list beats allow list; don’t-ask continues after a blocked step; usage policy; input/output scanners.
- Products: Chat vs Cowork vs Claude Code from weeks 1–3 — still in the Foundation story.

IF YOU ONLY DO FIVE THINGS

- 1) Skilljar Foundation/API course.
- 2) Official tool-use and MCP docs.
- 3) One timed mock for pace.
- 4) Narrate RAG without naming Atlas as mandatory.
- 5) Know plan vs accept-edits vs don’t-ask.

After class logistics

PDFs (including Agent SDK and guardrails README/PPT) were being copied into the cohort repo's PDF folder. If a file is missing, ask the facilitator rather than scraping a random GitHub. Ping questions in the class chat; he offered personal chat then realised privacy settings blocked it — use the shared channel or the repo.

Official documentation

- **Anthropic Academy (Skilljar)**

<https://anthropic.skilljar.com>

Official courses and exam registration.

- **Claude Docs home**

<https://docs.claude.com/en/home>

Start here the morning of the exam if a UI has moved.

- **Tool use overview**

<https://docs.claude.com/en/docs/agents-and-tools/tool-use/overview>

High-yield API domain.

- **MCP**

<https://docs.claude.com/en/docs/agents-and-tools/mcp>

High-yield connector domain.

- **Usage policy**

<https://www.anthropic.com/legal/usage-policy>

RAI / allowed-use questions.

- **Strengthen guardrails**

<https://docs.claude.com/en/docs/test-and-evaluate/strengthen-guardrails>

How Anthropic talks about testing agents.

- **Help: where is the API documentation**

<https://support.claude.com/en/articles/8114490-where-can-i-find-your-api-documentation>

If you get lost between console and docs.

Study notes distilled from the UST CCD Foundation masterclass. Prefer Anthropic's current docs over session anecdotes. Not official exam material.